

RAG-Enabled Chat Application

Technical Documentation & Deployment Guide

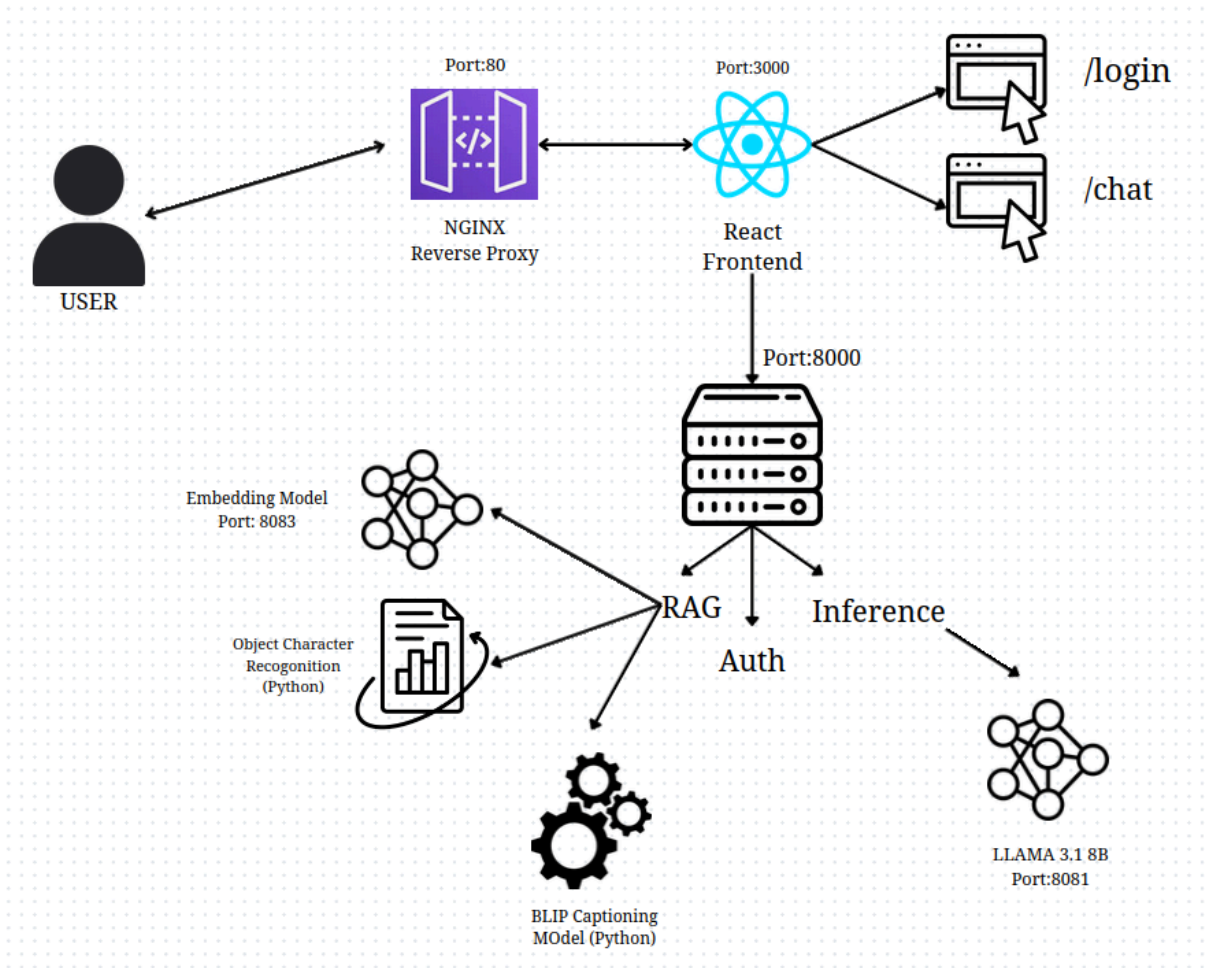


Table of Contents

1. System Overview
2. Technology Stack
3. Infrastructure Setup
4. Deployment Configuration
5. Application Architecture
6. API Documentation
7. System Internals

System Overview

This document provides comprehensive technical documentation for a Retrieval-Augmented Generation (RAG) enabled chat application. The system integrates modern web technologies with advanced language models to deliver intelligent conversational AI capabilities with document-based context awareness.

Technology Stack

Frontend Components

- **Framework:** React
- **Deployment:** Served via NGINX reverse proxy on port 3000

Backend Infrastructure

- **API Server:** FastAPI REST API
 - **RAG Engine:** Langchain for Retrieval-Augmented Generation
 - **Inference Server:** LLaMA.cpp with Meta LLaMA 3.1 8B model
 - **Embedding Model:** nomic-embed-text-v1 for document vectorization
 - **Database:** SQLite for user data persistence
 - **Vector Database:** FAISS for similarity search
-

Infrastructure Setup

NGINX Configuration

Installation

Navigate to the offline debian package directory and install NGINX:

```
sudo dpkg -i *.deb
```

Reverse Proxy Configuration

Create the NGINX configuration file:

```
sudo nano /etc/nginx/sites-available/reverse-proxy
```

Add the following configuration:

```
server {  
    listen 80;
```

```
server_name _;

# Proxy all requests to React frontend on port 3000
location / {
    proxy_pass http://localhost:3000;
    proxy_set_header Host $host;
    proxy_set_header X-Real-IP $remote_addr;
    proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
    proxy_set_header X-Forwarded-Proto $scheme;
}
}
```

Service Configuration

Remove default configuration and enable the reverse proxy:

```
sudo rm /etc/nginx/sites-enabled/default
sudo ln -s /etc/nginx/sites-available/reverse-proxy /etc/nginx/sites-enabled/
```

Test and start the service:

```
sudo nginx -t
sudo systemctl daemon-reload
sudo systemctl enable nginx
sudo systemctl start nginx
```

Deployment Configuration

LLaMA.cpp Server Setup

Environment Configuration

Configure the library path for LLaMA.cpp:

```
export LD_LIBRARY_PATH=build/ggml/src:build/src:$LD_LIBRARY_PATH
```

Model Server Deployment

Launch the LLaMA inference server:

```
./build/bin/llama-server \
-m "/path/to/model/Meta-Llama-3.1-8B-Instruct-Q4_K_M.gguf" \
-np -c 2048 --port 8081 --host 0.0.0.0 \
--chat-template llama3
```

Launch the embedding server:

```
./build/bin/llama-server \  
-m "/path/to/model/nomic-embed-text-v1.Q4_K_M.gguf" \  
-c 2048 --port 8083 --host 0.0.0.0 \  
--embeddings
```

Automated Service Management

Startup Script Creation

Create an executable startup script:

```
nano startup_services.sh  
chmod +x startup_services.sh
```

Startup Script Content

```
#!/bin/bash  
  
# Ubuntu startup script for launching all services  
# Ensure paths are updated according to your deployment setup  
  
LOG_DIR="$HOME/logs"  
mkdir -p "$LOG_DIR"  
  
echo "Initializing all services..."  
  
# Backend Server  
echo "Starting backend server..."  
cd /path/to/backend && python3 main.py > "$LOG_DIR/backend.log" 2>&1 &  
BACKEND_PID=$!  
echo "Backend server started with PID: $BACKEND_PID"  
  
# Frontend Server  
echo "Starting frontend server..."  
cd /path/to/frontend && PORT=3000 HOST=ip npx serve build -s > "$LOG_DIR/frontend.log"  
2>&1 &  
FRONTEND_PID=$!  
echo "Frontend server started with PID: $FRONTEND_PID"  
  
# LLaMA Inference Server  
echo "Starting LLM inference server..."  
cd /path/to/llama.cpp && [inference_command] > "$LOG_DIR/llama.log" 2>&1 &  
LLM1_PID=$!
```

```
echo "LLaMA inference server started with PID: $LLM1_PID"

# Embedding Server
echo "Starting embedding server..."
cd /path/to/llama.cpp && [embedding_command] > "$LOG_DIR/embed.log" 2>&1 &
LLM2_PID=$!
echo "Embedding server started with PID: $LLM2_PID"

# Process ID Management
echo "$BACKEND_PID" > "$LOG_DIR/pids.txt"
echo "$FRONTEND_PID" >> "$LOG_DIR/pids.txt"
echo "$LLM1_PID" >> "$LOG_DIR/pids.txt"
echo "$LLM2_PID" >> "$LOG_DIR/pids.txt"

echo "All services initialized successfully!"
echo "Log files available in: $LOG_DIR"
echo "Process IDs saved to: $LOG_DIR/pids.txt"
```

SystemD Service Configuration

Create a system service for automated startup:

```
sudo nano /etc/systemd/system/my-services.service
```

Service configuration:

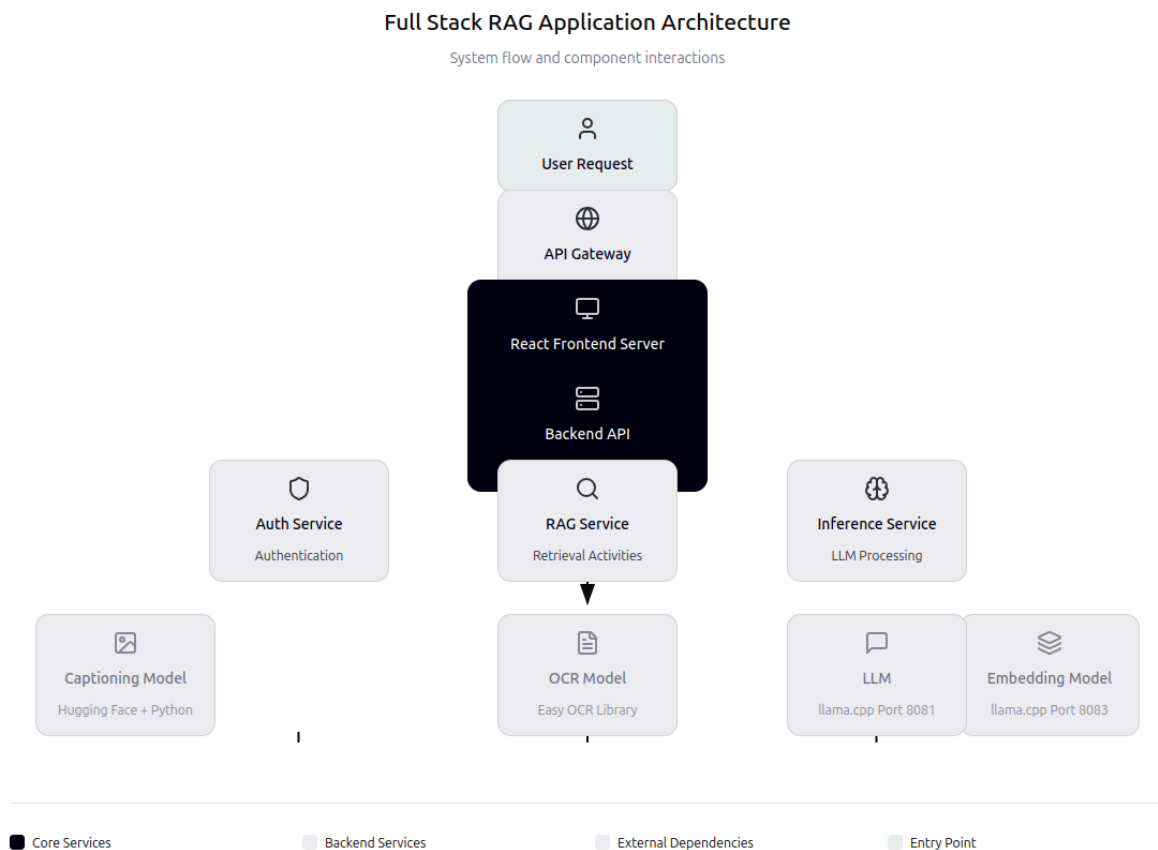
```
[Unit]
Description=RAG Chat Application Services
After=network.target graphical-session.target
Wants=network.target

[Service]
Type=forking
User=your_username
Group=your_username
WorkingDirectory=/home/your_username
ExecStart=/home/your_username/startup_services.sh
ExecStop=/home/your_username/stop_services.sh
Restart=on-failure
RestartSec=5
Environment=HOME=/home/your_username

[Install]
WantedBy=multi-user.target
```

Enable and start the service:

```
sudo systemctl daemon-reload
sudo systemctl enable my-services.service
sudo systemctl start my-services.service
```



Application Architecture

Frontend Components

Authentication Interface

- **Login Page:** Secure user authentication with username/password validation
- **Registration System:** New user account creation with availability checking

Chat Interface Features

- **Session Management:** Clear chat functionality with frontend-only history storage
- **Theme Toggle:** Light/dark mode switching capability
- **Logout Functionality:** Secure session termination with backend cleanup
- **Document Upload:** RAG-enabled file processing with FAISS indexing
- **Real-time Messaging:** Streaming inference with abort capability

Backend Architecture

Service Distribution

- **Frontend:** React application on port 3000 (NGINX reverse proxy from port 80)
 - **Backend API:** FastAPI server on port 8000 for authentication and RAG processing
 - **LLaMA Inference:** Model server on port 8081 for text generation
 - **Embedding Service:** Dedicated server on port 8083 for document vectorization
-

API Documentation

Authentication Endpoints

User Registration

POST /register

Request Body:

```
{  
  "username": "string",  
  "password": "string"  
}
```

Responses:

- **200 OK:** Registration successful
- **409 Conflict:** Username already exists

User Authentication

POST /login

Request Body:

```
{  
  "username": "string",  
  "password": "string"  
}
```

Success Response:

```
{  
  "message": "Login successful",  
}
```

```
"username": "string"
}
```

Error Response:

- **401 Unauthorized:** Invalid credentials

Session Termination

POST /logout

Request Body:

```
{
  "username": "string"
}
```

Response:

```
{
  "message": "Logged out successfully."
}
```

Chat and Inference Endpoints

Streaming Message Processing

POST /message/stream

Request Body:

```
{
  "chatHistory": "User: Hi\nAssistant: Hello! How can I help you?",
  "message": "What is Retrieval-Augmented Generation?",
  "model": "llama3.1:8b",
  "ragStatus": true,
  "session_id": "<uuid>",
  "user": "<username>"
}
```

Functionality:

- Processes user queries with optional RAG context
- Performs FAISS similarity search when **ragStatus = true**
- Includes top 4 relevant document chunks in system prompt

- Returns streaming response using `text/event-stream` format

RAG Session Management

Session Creation

POST `/rag/create-session`

Response:

```
{
  "session_id": "<uuid>",
  "status": "success",
  "message": "RAG session created successfully"
}
```

Session Deletion

DELETE `/rag/session/{session_id}`

Session Status

GET `/rag/session/{session_id}/status`

Table Extraction

GET `/rag/session/{session_id}/tables`

Active Sessions

GET `/rag/sessions`

Document Processing

Document Upload and Indexing

POST `/rag/upload/{session_id}`

Form Data:

- `file`: Document file (PDF, DOCX, TXT, PNG, JPG, JPEG)
- `include_tables`: Boolean (default: true)

Processing Pipeline:

1. Text extraction using pdfplumber, docx, and plain text parsers
2. Table extraction and restructuring

3. OCR and image captioning using BLIP + EasyOCR
4. Document chunking via Langchain RecursiveCharacterTextSplitter
5. Embedding generation using LLaMA.cpp model
6. FAISS indexing with Inner Product similarity

Response:

```
{
  "status": "success",
  "message": "Document processed successfully",
  "chunks_created": 24,
  "processing_time": 4.73,
  "filename": "example.pdf",
  "session_id": "<uuid>"
}
```

Document Removal

DELETE /rag/document/{session_id}

System Internals

Document Chunking Strategy

Configuration Parameters:

- **chunk_size:** 800 characters
- **chunk_overlap:** 100 characters
- **Splitting Logic:** Intelligent text segmentation using newlines, punctuation, and whitespace

Embedding Generation

Model Configuration:

- **Server:** LLaMA.cpp on port 8083
- **Model:** nomic-embed-text-v1
- **Preprocessing:** Text prefixed with "**search_query:**" for schema compliance

Similarity Search Implementation

FAISS Configuration:

- **Index Type:** IndexFlatIP with L2-normalized vectors
- **Retrieval Count:** max(4, total_chunks)

- **Similarity Threshold:** 0.75 minimum

Prompt Template System

Message Format Conversion: OpenAI-style messages converted to LLaMA chat format:

```
<|begin_of_text|><|start_header_id|>role<|end_header_id|>
```

```
content<|eot_id|>
```

Session Lifecycle Management

Memory Management:

- All sessions maintained in-memory
- Each session contains document data, metadata, chunks, and vector index
- **Idle Timeout:** 2 hours
- **Cleanup Schedule:** Background task every 1 hour

Conclusion

This documentation provides a comprehensive guide for deploying and maintaining the RAG-enabled chat application. The system architecture ensures scalability, security, and optimal performance for document-enhanced conversational AI interactions.

For additional support or configuration assistance, please refer to the respective technology documentation or contact the development team.